

# Monary

Really Fast Analysis with NumPy and MongoDB

*Matt Cotter & Kyle Suarez*

# Imagine that you are a...

...data scientist

...biologist

...MongoDB user

...someone that has taken a taxi

**...and you want to analyze data.**

# What if you only need an average?

If the data is stored in a...

- ❖ List of Python Dicts

~12 million numbers/sec : (

- ❖ Python List

110 million numbers/sec : |

- ❖ `numpy.ndarray`

500 million numbers/sec : )

# NumPy

- ❖ Numerical Python
- ❖ Arrays and Matrices
- ❖ Lots of math

How do we get data from our database into  
NumPy?

*Monary*

# About us

Kyle Suarez

 @ksuarz

 ksuarz

Matt Cotter

 @MattWCotter

 machyne

College seniors  
Interns at MongoDB Inc.  
Core contributors to Monary

# Why MongoDB?

# MongoDB: a NoSQL database

- ❖ open-source document database
- ❖ data in binary JSON format

```
{  
  "names": ["Matt", "Kyle"],  
  "ages": [21, 21],  
  "talk": {"pygotham": "Monary"}  
}
```



# Distribution

- ❖ Scale
- ❖ Read Capacity

# Data is unpredictable.

## ❖ Flexible schemas:

- Accommodates varying data structures
- Not all documents need to have the same fields

## ❖ Works well with NumPy's masked arrays


# What happens in between?

MongoDB

????

NumPy

# PyMongo Workflow



# **This is pretty fast:**

With the server and processes both running on this laptop...

About 150,000 documents read per second.


# But...

- ❖ Object creation overhead is expensive.
- ❖ Dicts are only intermediate storage vessels.

## Can we get faster?

A black and white photograph of Bruce Lee in a dynamic martial arts pose, likely from the movie 'Enter the Dragon'. He is shirtless, with his hands raised in a defensive or offensive stance. The background shows a blurred outdoor setting with architectural elements. A semi-transparent horizontal band across the middle of the image contains the text 'Enter Monary.' in a bold, black, sans-serif font.

**Enter Monary.**



## Monary Workflow





**This is *really* fast!**

Same test as before...

This time over 1,700,000 reads per second.


**What is Monary?**

**Monary:** the bridge between



mongoDB

&



NumPy

# Monary Overview

- ❖ author: David J. C. Beach
- ❖ operations: CRUD only
- ❖ dependencies: NumPy
- ❖ driver: MongoDB C Driver 0.98.0

# How does it work?

- ❖ Implemented in C:
  - MongoDB C driver
  - MongoDB C BSON library
- ❖ NumPy + ctypes:
  - Create NumPy arrays
  - Pass pointers to the arrays into C

**Demo**

## Table Of Contents

- Monary 0.3.0 Documentation
  - Overview
  - Dependencies
  - Issues
  - Contributing
  - Changes
  - About This Documentation
  - Indices and tables

## Next topic

[Installing / Upgrading](#)

## This Page

[Show Source](#)

## Quick search

Go

Enter search terms or a module, class or function name.

# Monary 0.3.0 Documentation

## Overview

**Monary** provides a Python interface for fast column queries from [MongoDB](#). It is recommended over PyMongo when requiring large bulk reads from a database into [NumPy](#) arrays.

Note that Monary is still in beta. There are no guarantees of API stability; furthermore, dependencies may change in the future.

Monary is written by [David J. C. Beach](#).

### *Installing / Upgrading*

Instructions on how to get the distribution.

### *Tutorial*

Getting started quickly with Monary.

### *Examples*

Examples of how to perform specific tasks.

### *Type Reference*

In-depth explanation of how Monary handles BSON types.

### *Frequently Asked Questions*

Frequently asked questions about Monary.

[See tinyurl.com/monary-docs](http://tinyurl.com/monary-docs)

## Dependencies

Monary depends on [PyMongo](#) and [NumPy](#).

Monary uses the [MongoDB C driver](#), which comes bundled as part of the module.

## Issues

All issues can be reported by opening up an issue on the Monary [BitBucket issues page](#).

## Contributing

Monary is an open-source project and is hosted on [BitBucket](#). To contribute, fork the project and send a pull request.

See the [Contributors](#) page for a list of people who have contributed to developing Monary.

# Now install it!

On PyPI:

[pypi.python.org/pypi/Monary](https://pypi.python.org/pypi/Monary)

With pip:

```
$ pip install monary
```



# The Future of Monary

# Potential new features

- ❖ Release all our features
- ❖ Pandas integration
- ❖ Bulk updates, upserts
- ❖ Array extraction
- ❖ Logging

# Contributing

Author: David J. C. Beach

License: Apache License Version 2.0

Open source on BitBucket:

[bitbucket.org/djcbeach/monary](https://bitbucket.org/djcbeach/monary)

# Thanks!

Our mentors at MongoDB:

- ❖ A. Jesse Jiryu Davis
- ❖ Jason Carey

The author of Monary:

- ❖ David J. C. Beach

# Questions?

**Source:** [bitbucket.org/djcbeach/monary](https://bitbucket.org/djcbeach/monary)

**PyPI:** [pypi.python.org/pypi/Monary](https://pypi.python.org/pypi/Monary)

**Documentation:** [tinyurl.com/monary-docs](https://tinyurl.com/monary-docs)

**This Talk:** [github.com/Machyne/pygotham-monary](https://github.com/Machyne/pygotham-monary)